Engineering Assignment

Congratulations on being shortlisted!

We believe the best way to evaluate an engineer is to let them *build*. So, below are two projects. You need to build **one of them**.

> **You are allowed (and encouraged) to use AI.** We are interested in knowing about the approach you take and decisions you make, not your syntactic fluency.

Before You Begin

- **Attention to detail** and **completeness** of your solution will be appreciated.
- **Make it runnable.** Another developer should be able to clone your repo and get it running with minimal friction. Provide necessary steps to make it so, in your preferred format.
- **Follow best practices** for whatever stack you choose. We're not prescribing a tech stack, pick whatever you think is the best. No limitations on platform, it’s up to you.
- Commit your work regularly. Every project is built through iterations, we want to see your work come to life.
- **(Optional) Ship it.** Deploy it somewhere. A live link goes a long way.

Option A: Voice-First Task App

> *Build a task management app where the primary input method is **natural language voice**.*

A user should be able to speak something like *"Remind me to submit the quarterly report by next Friday"* — and your app

should figure out the rest.

What we expect

1. Voice input that extracts structured task data automatically
    - Each Task can have the following fields (or more):
        - title
        - description
        - due date
2. Tasks can be **created**, **completed**, **cancelled**, and **delayed**
3. Users should be authenticated
4. An analytics dashboard with visual graphs covering:
    - Tasks completed on time
    - Tasks currently pending
    - Tasks that were delayed

What we're evaluating

- How you handle ambiguity in natural language.
- How you think about UI/UX.
- How useful the KPIs you have chosen for the analytics page would be in practice.
- Your design decisions in case of ambiguous requirements.

Option B: Legal Research RAG Chatbot

> Build a Retrieval-Augmented Generation chatbot over the **Bharatiya Nyaya Sanhita (BNS)**

The source documents are available here: Link to documents

What we expect

1. A proper RAG pipeline, **not** "upload PDF to ChatGPT and call it a day"
2. Vector embeddings with an **open-source vector database running locally**
3. An **open-weight LLM** for generation (it doesn't need to run locally, but the model must be open-weight)
4. It should be **fast**.

What we're evaluating

- Your chunking strategy. Your retrieval quality.
- How you think about UI/UX.
- How you handle questions that span multiple sections of the law.
- Whether your system actually gives *accurate* answers. Bonus points if you've written evals or tests to prove it.

Submission

- Push your code to a **public GitHub repository**
- Include a README with setup instructions (or anything else you might want to add)
- Share the repository link with your recruiter

Deadline: You'll have **3 days** from receiving this assignment.

> *“I’m gonna keep getting better, keep getting better, and one day, everything's gonna pay off.”* - Giannis Antetokounmpo

Good luck! :)